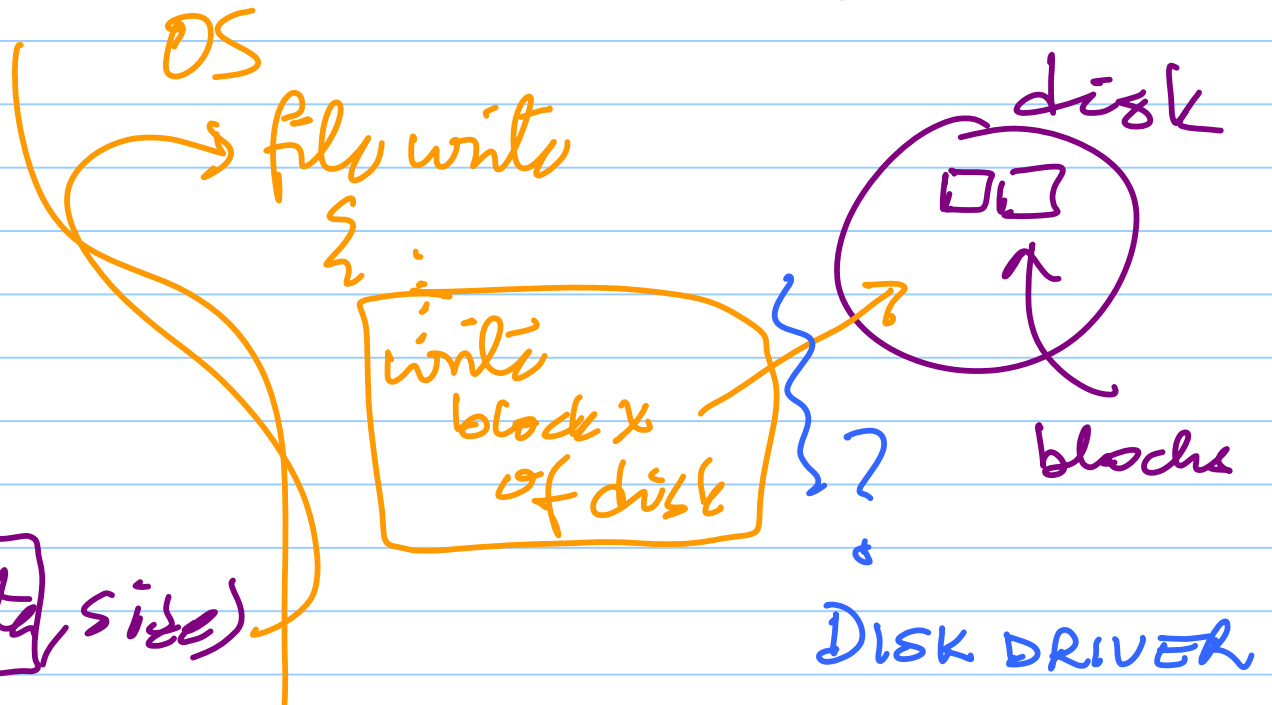


Application

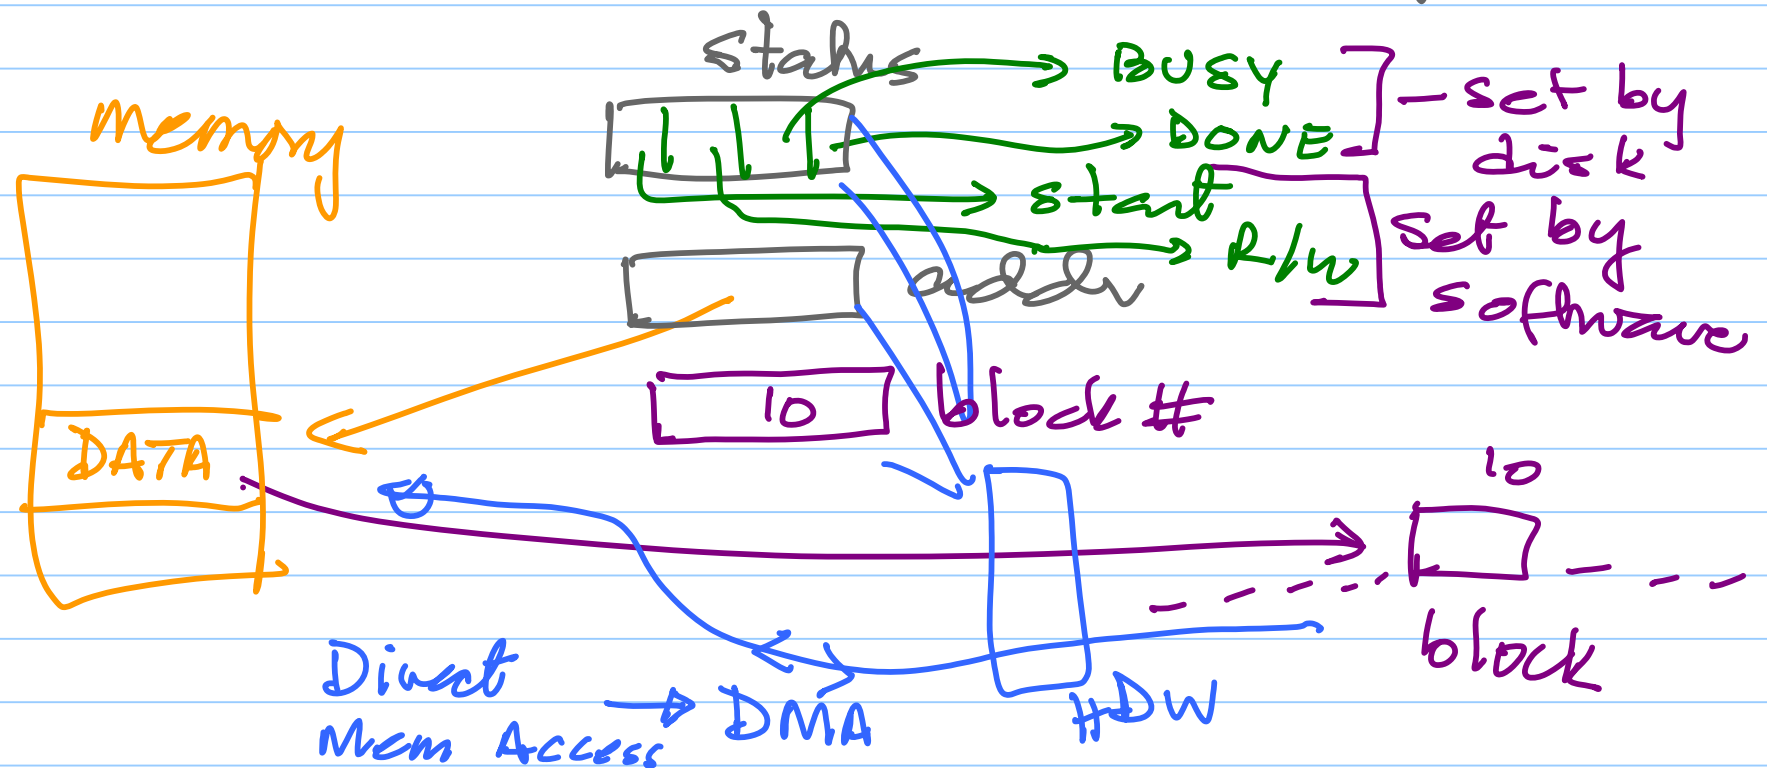
∴ program
read/write
from/to
files

write (fd, data, size)

DISK I/O



Disk interface \rightarrow set of registers



// polled I/O

write a block to disk (b#, addr)

{ while (busy); // wait

load b# → b# register

load addr → addr reg

Set r/w bit to w

set start bit

while (not done); // wait for

write
to complete

}

Atomic

polling

Interrupt driven I/O.

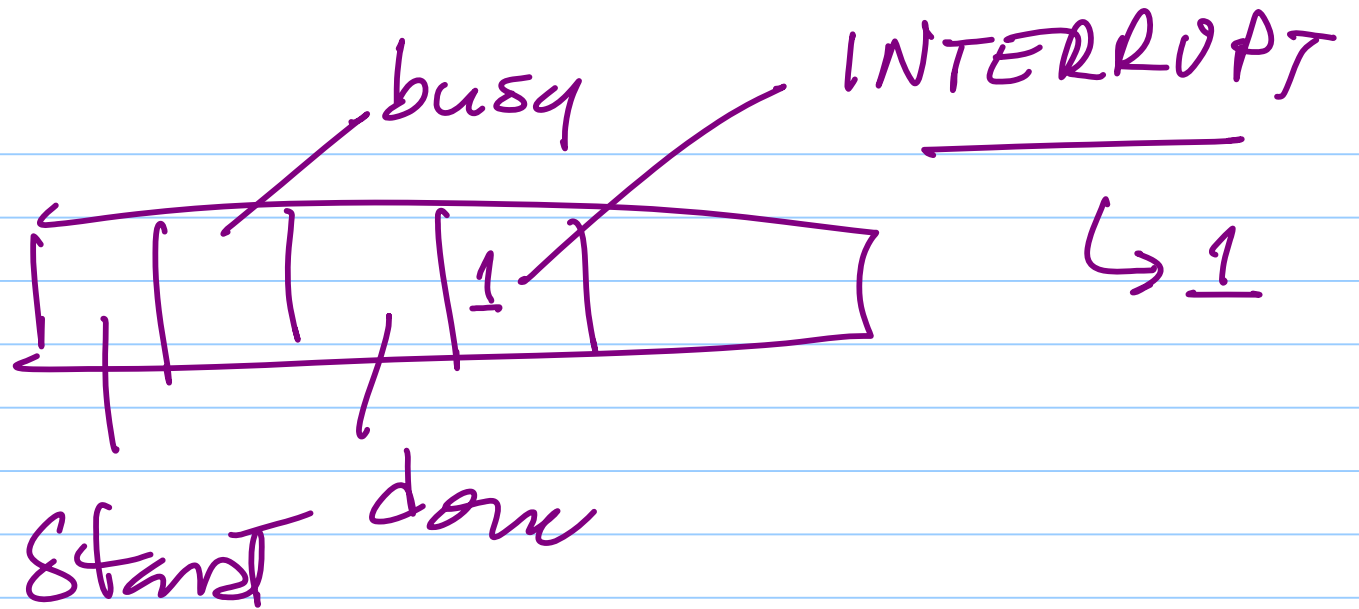
↳ driver starts things up

→ blocks

→ an interrupt
wakes it up

→ hardware
signal
generated
at the
end of
an operation

Status
reg



INT bit is 1

write ... (b#, addr)

{ P(disksem) $\xrightarrow{\text{set to 1}}$ q of processes waiting
Set addr, b#, w bit, start for I/O

\rightarrow P(disk-int)

$\xleftarrow{\text{set to 0}}$
V(disksem)

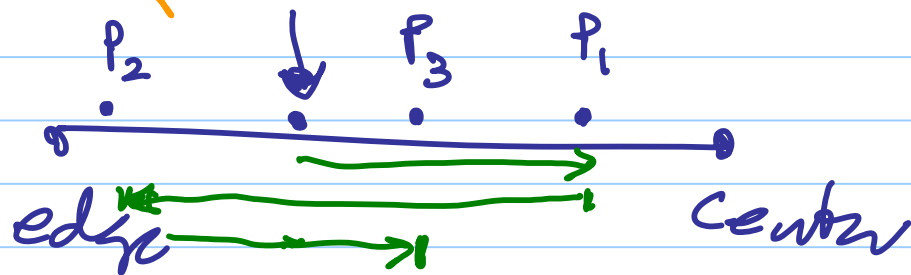
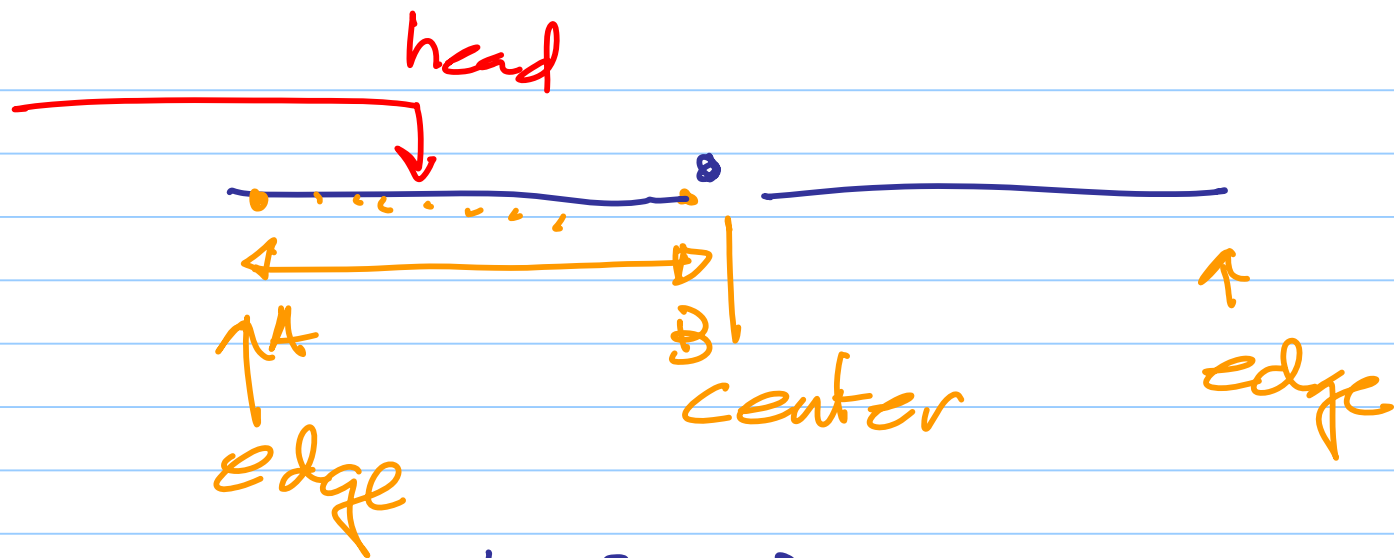
FCFS I/O scheduling

diskinthandler()
{ V(disk-int) }

DISK | scheduling
↓
head

→ only for mechanical disks

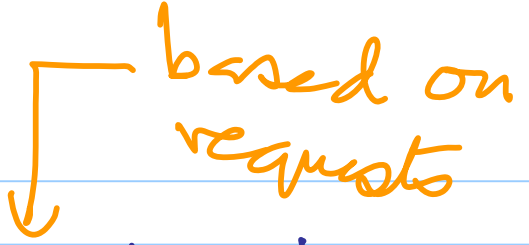
→ what order should waiting I/O processes be using the disk



FCFS

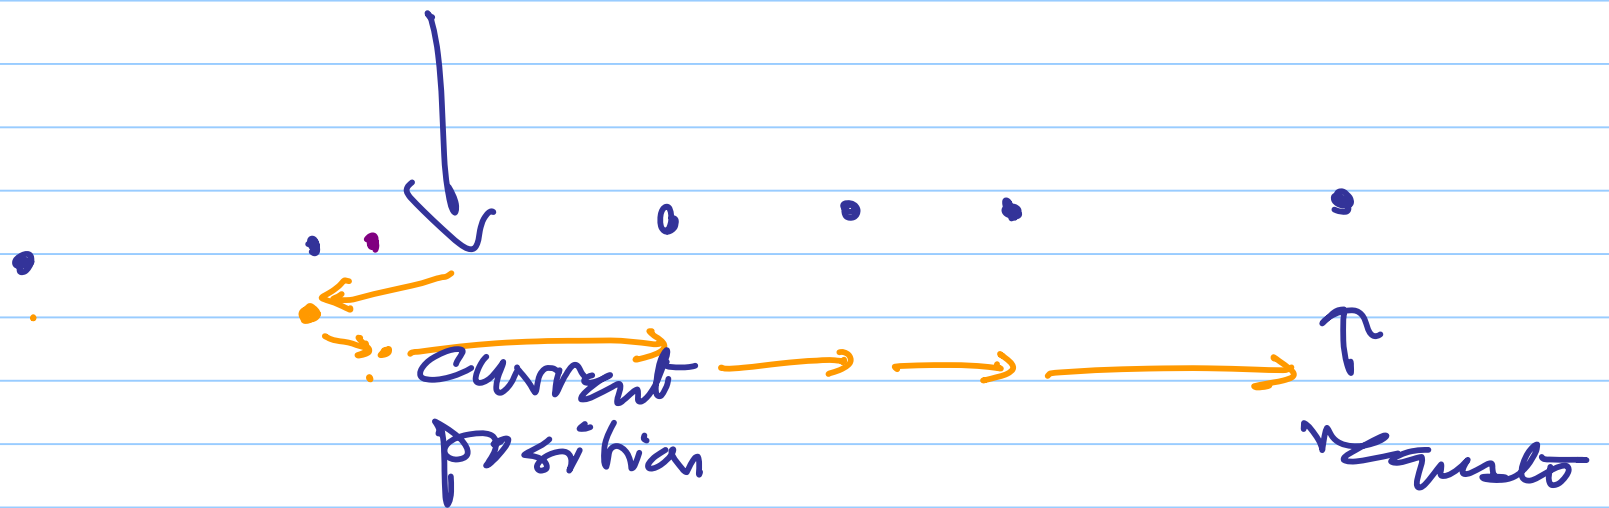
→ head moves "randomly" all over the disk

based on requests



→ may have too much head movement

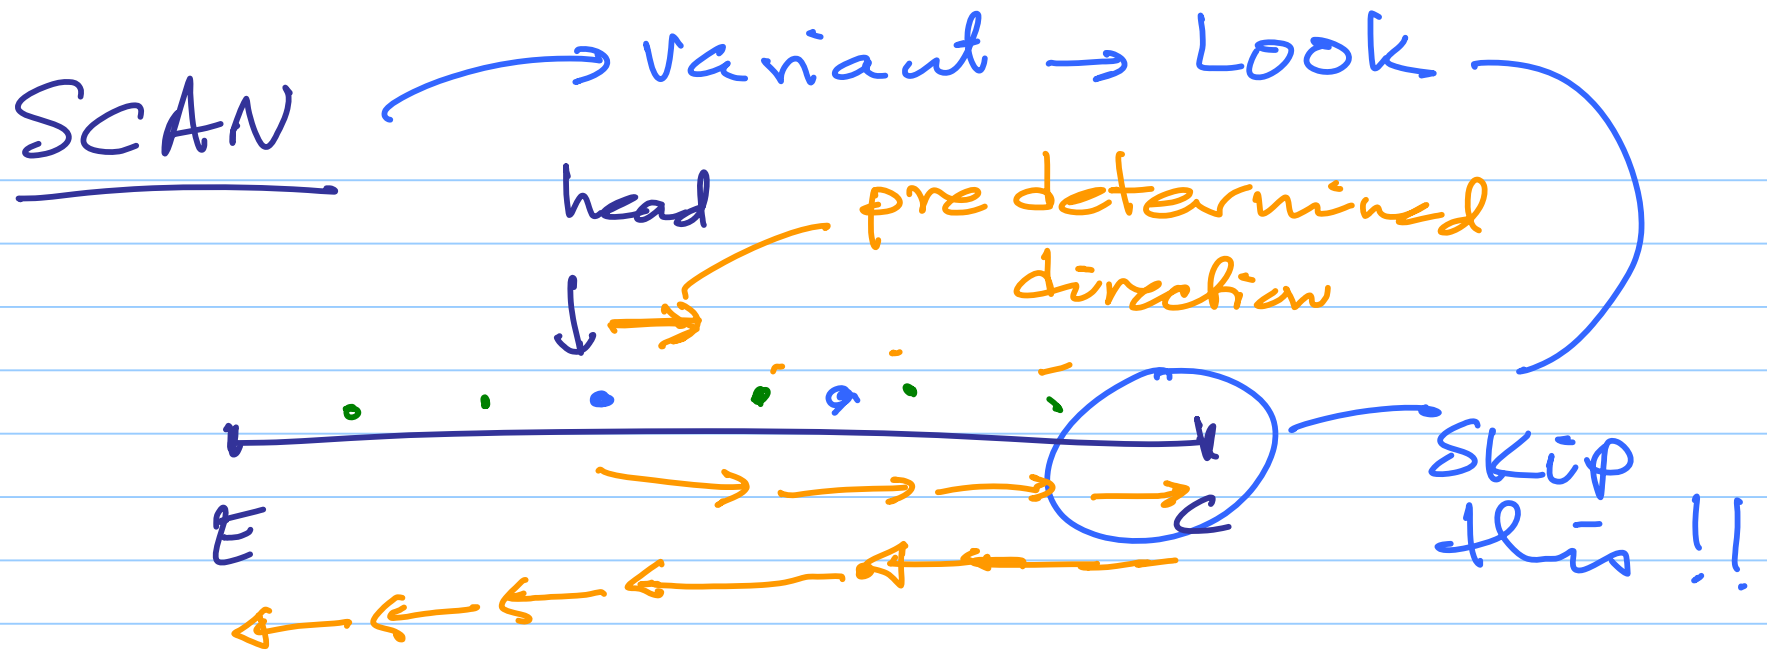
SSTF \rightarrow shortest seek time first



\rightarrow low head movement, but starvation

→ Under low load — disk access is
FCFS



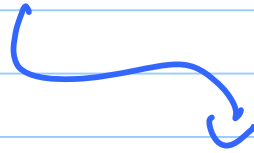


Travel Edge to center & then
center to edge

SCAN → LOOK



C-SCAN

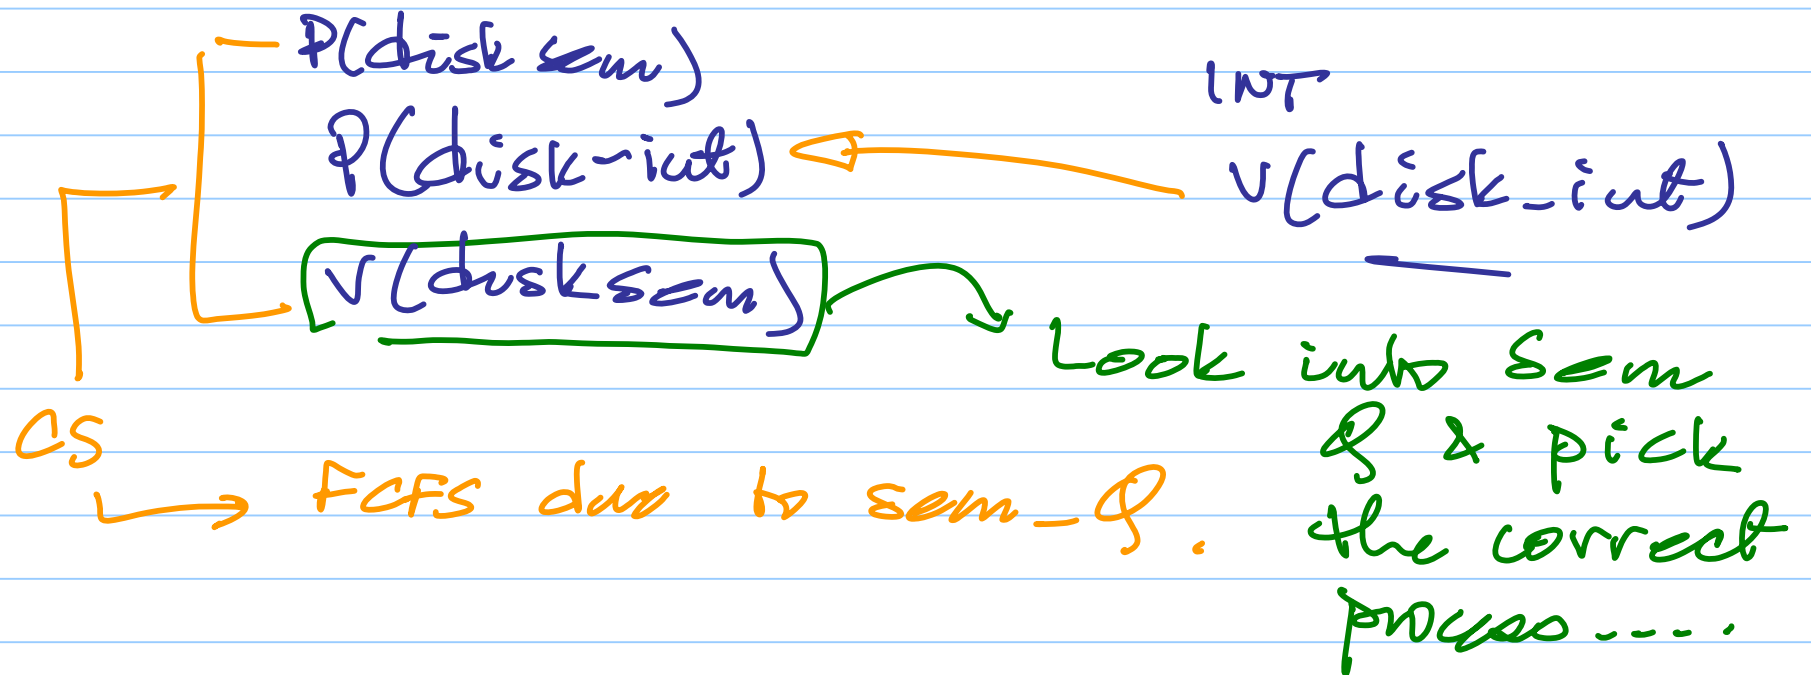


C-LOOK
↪

↳ circular scan

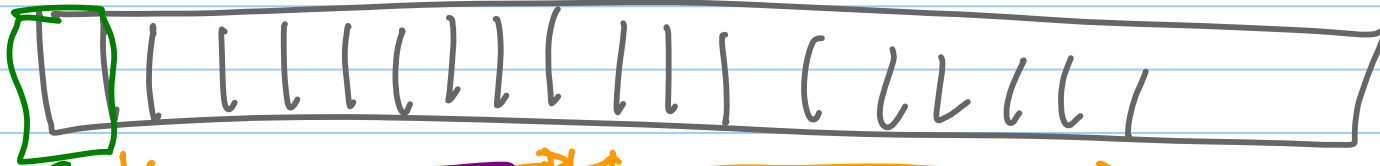
- edge to center
- retract
- edge to center

How to do non FCFS ?



Disk Management

blocks



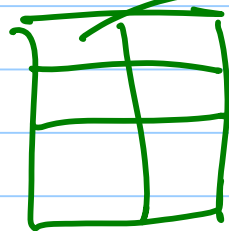
Part 1

Part 2

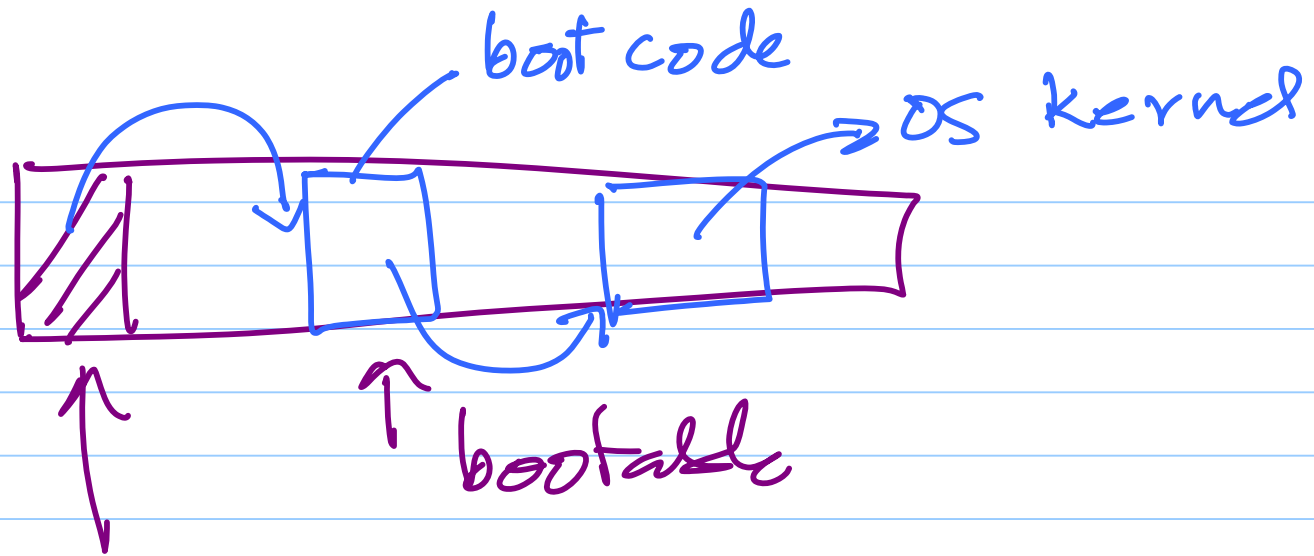
① partitioning

bootable partition

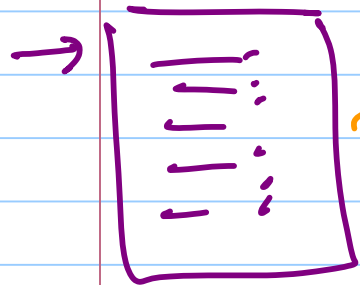
partition table



partition #, name, start block
end block (size) ...



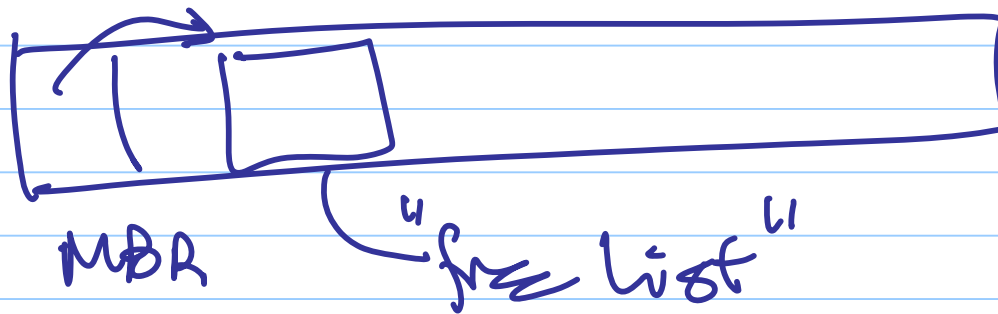
MBR → master boot record.



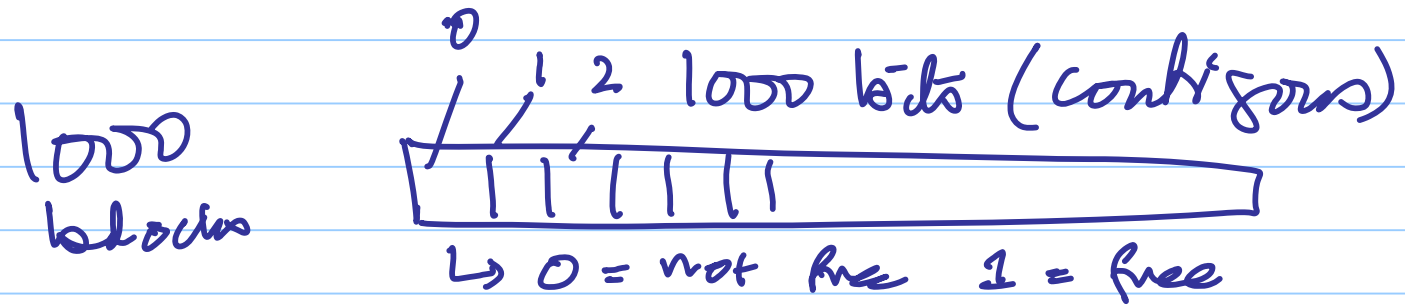
boot rom

find partition
find MBR
find boot-code

free disk blocks partition

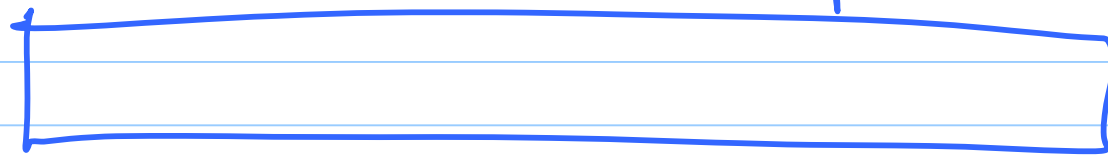


↳ bit vector



Swap space

Swap partition



- a swap file → a file (large)
- a " partition → full partition
- a " SPACE → contiguous area in a partition